


Structured Encoder for Ancient Logograms


EMANUELE MESSINA
SALVATORE CIMMINO

2025

Evolution of Chinese Scripts




Straightforward Approach



Radicals Examples

Radical	water		fire	
Script	Regular	Seal	Regular	Seal
	水	𣎵	火	𤇀
Characters	lake	river	hot	burn
Regular Script	湖	江	熱	燒
Seal Script	𣎵胡	𣎵工	𤇀𤇀	𤇀𤇀

Hierarchical Approach



© 2006 The Authors



Characters selection





Image generation






Splits

- How many images? → Coupon collector's problem, rounded up
- 90% Train (35k) / 10% test (4k)
- Labels format:
 - SEAL: `char x1 y1 x2 y2 radical`
 - YOLO: `char_id cx cy bw bh`
- Calculated image mean and std over the train split

Class weights

- High (intrinsic) imbalance in radicals' distribution
- Need superclass weights (weighted CE)
- Also did for subclasses (although more uniform)

$$w_c = \frac{\sum_{i \in C} n_i}{|C| n_c}$$




Model


Building SEAL

```
187 (cls_score_dummy): Linear(in_features=1024, out_features=2, bias=True)
188 (bbox_pred): Linear(in_features=1024, out_features=8, bias=True)
189 (superclassifier_funnel): Sequential(
190   (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
191   (1): ReLU()
192   (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
193   (3): ReLU()
194   (4): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
195   (5): ReLU()
196   (6): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1))
197   (7): ReLU()
198   (8): Conv2d(256, 256, kernel_size=(2, 2), stride=(1, 1), padding=(1, 1))
199   (9): ReLU()
200   (10): Conv2d(256, 256, kernel_size=(2, 2), stride=(1, 1), padding=(1, 1))
201   (11): ReLU()
202   (12): AdaptiveAvgPool2d(output_size=3)
203 )
204 (superclassifier_head): Sequential(
205   (0): Flatten(start_dim=1, end_dim=-1)
206   (1): Linear(in_features=2304, out_features=1024, bias=True)
207   (2): ReLU()
208   (3): Linear(in_features=1024, out_features=214, bias=True)
209 )
210 (subclassifier_funnel): Sequential(
211   (0): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
212   (1): ReLU()
213   (2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
214   (3): ReLU()
215   (4): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
216   (5): ReLU()
217   (6): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1))
218   (7): ReLU()
219   (8): Conv2d(1024, 1024, kernel_size=(2, 2), stride=(1, 1), padding=(1, 1))
220   (9): ReLU()
221   (10): Conv2d(1024, 1024, kernel_size=(2, 2), stride=(1, 1), padding=(1, 1))
222   (11): ReLU()
223   (12): AdaptiveAvgPool2d(output_size=3)
224 )
225 (subclassifier_junction): Sequential(
226   (0): Conv2d(1280, 1280, kernel_size=(1, 1), stride=(1, 1))
227   (1): ReLU()
228   (2): Conv2d(1280, 1280, kernel_size=(1, 1), stride=(1, 1))
229   (3): ReLU()
230   (4): Flatten(start_dim=1, end_dim=-1)
231   (5): Linear(in_features=11520, out_features=2048, bias=True)
232   (6): ReLU()
233 )
234 (sub_classifiers): ModuleList(
235   (0): Linear(in_features=2048, out_features=328, bias=True)
```


Final Architecture




Final Architecture




ResNet-50 FPN




Final Architecture




Final Architecture




Final Architecture




Final Architecture




Final Architecture




Final Architecture




Final Architecture



Final Architecture



Final Architecture




Loss


Box Regression
+ Binary Classification Loss

Radical
Classification Loss


Character
Classification Loss

$$L = L_f + \lambda_r L_r + \lambda_c L_c$$



Faster R-CNN Forward Pass



Custom Forward Pass



Experiments



Early experiments

- Fully linear, big latent space dimension (8k), 3GB, ROI size 7
- Provisory 725 and 2k datasets (pruned with Classical Corpus frequency list)
- Showed the need of distancers otherwise gradient/task conflicts with regression loss
- Converged on provisory datasets without Hierarchical Classification
- Only Hierarchical version converged on Full Dataset
- Moved to Convolutional Model: faster, slimmer (145M params, 1GB), more accurate (+10% mAP)
- ROI size increased to 8
- Anchors: {50, 70} with aspect ratios {1:1 , 2:1}



Training


- Adam (per param momentum, fast, less sensitive)? Exploding/unstable loss
- Gradient clipping (max norm 10) after that
- SGD (default values: momentum 0.9, weight decay $1e-5$), slow but stable
- Initial Learning Rate $1e-3$ + StepLR (gamma 0.8) every epoch (steadiest descent config)
- Batch size 2, Epochs 6
- $\lambda_r = 0.99$, $\lambda_c = 0.95$

Results



YOLO11x (57M, 111MB)

TARGETS




PREDICTIONS




SEAL (Hierarchical)

RADICALS



CHARACTERS




METRICS


mAP@0.75	0.57
mAR	0.65

Confusion Matrices (Non-Hierarchical)


RADICALS



CHARACTERS



Qualitative inspection



Future Work

- Increase ROI Size
- Increase Backbone Channels
- More dataset augmentations:
 - colors
 - backgrounds
 - deteriorations
 - negative examples
- Junction layer ablation

